


New observations improve prediction of recurrent updates

State history + observation beats the state-history-only fit in all 15 policies per protocol.



Each colored line: one policy; 3 training conditions × 5 learner seeds. Models use held-out trajectories and training-only state PCs. Errors score only newly predicted current-state coordinates. Targeted collection covers the first 4 rounds; predictive fit is not causal identification.