Cumulative episode return across rounds · Shared 24-profile familiar evaluation

All models: 24 training profiles, 2,400 episodes each per protocol, including the single-partner-trained RNN.

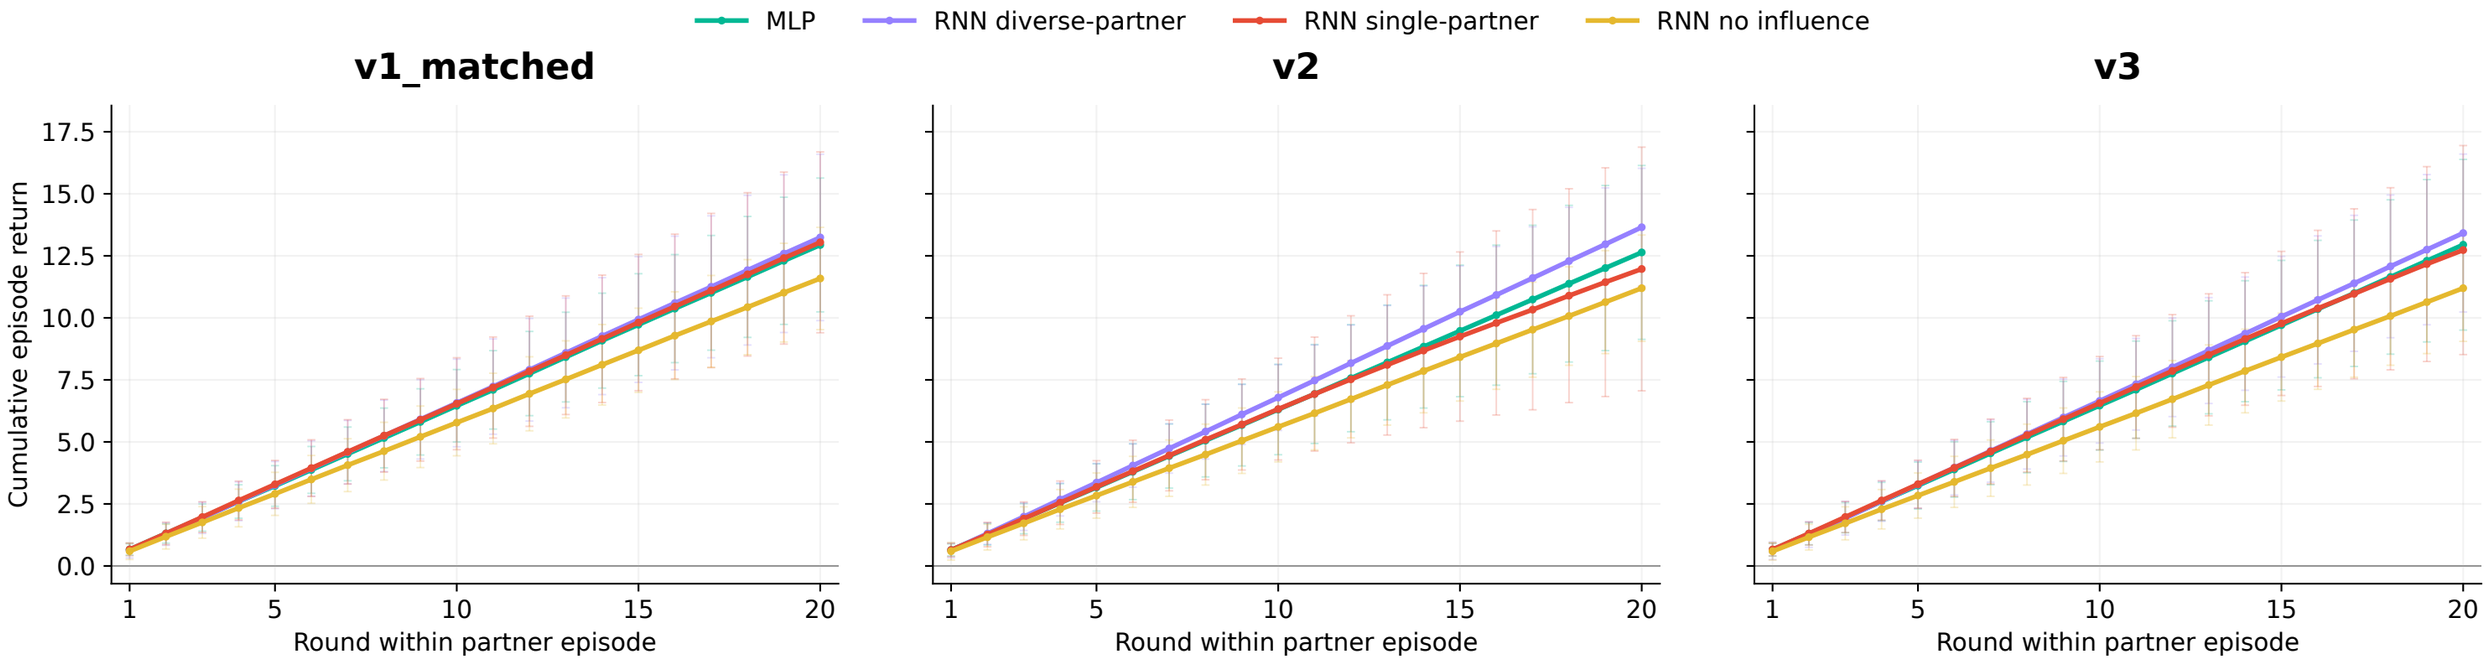

Means weight episodes equally across five learner seeds. Error bars: ±1 pooled episode sample SD, including differences between seeds.
Saved post-training evaluations; each episode has 20 rounds. Cumulative SD is computed from episode prefixes, preserving correlations between rounds.