Final assignment matches partner's faster task — selected training seeds

v1_matched | *v2* | *v3*

MLP — Seed 1 | Seed 4 | Seed 4

RNN — Seed 5 | Seed 2 | Seed 3

RNN single-partner — Seed 5 | Seed 4 | Seed 5

RNN no influence — Seed 3 | Seed 5 | Seed 5

Partner on faster task (%)

Round within partner episode

One policy per experiment and condition, selected by final logged training return among seeds 1–5.
Each curve averages 440 held-out episodes. Seed choices match the best-of-five return figure.