Current observation x_t faster-goal decoding · v1_matched, v2, v3 · direct input, fixed t0
Current observation x_t at every environment step t=0–400; no temporal averaging; five policy seeds; shading: 95% bootstrap CI.
Current observation x_t (249 inputs) · Fit: 24 profiles ×16 episodes; test: all90 ×4; five policy seeds.