Updated memory h_t faster-goal decoding · v1_matched, v2, v3 · direct input, fixed t0
Updated memory h_t at every environment step t=0–400; no temporal averaging; five policy seeds; shading: 95% bootstrap CI.
Updated memory h_t (128 inputs) · Fit: 24 profiles ×16 episodes; test: all90 ×4; five policy seeds.