Single-state faster-goal decoding · v1_matched, v2, v3 · direct input, fixed t0
h_t at every environment step t=0–400; no temporal averaging; five policy seeds; shading: 95% bootstrap CI.
Fit: 24 training profiles × 16 episodes; test: all 90 profiles × 4 episodes; five seeds.
Single-partner control omitted: one-class training support.